

STAT 243 PS 7

Junyuan Gao(SID:26484653)

November 17, 2017

1 Q1

Suppose the estimates of coefficients of 1000 datasets are $\hat{\beta}_1, \dots, \hat{\beta}_{1000}$. We can calculate the standard error of $\hat{\beta}_1, \dots, \hat{\beta}_{1000}$ and compare it to the mean of $se(\hat{\beta}_1), \dots, se(\hat{\beta}_{1000})$. If the two values are similar, we can say that the standard error properly characterizes the uncertainty of the estimated regression coefficient.

2 Q2

$$\|A\|_2 = \sup_{\|z\|_2=1} \sqrt{(Az)^T(Az)} = \sup_{\|z\|_2=1} \sqrt{z^T A^T A z}$$

Since $A^T A$ is symmetric, we can write eigen decomposition of $A^T A$ as

$$A^T A = U^T \Sigma U = U^T \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} U$$

where λ_i 's are the eigenvalues of A .

Plug it into $\|A\|_2$, we can get

$$\|A\|_2 = \sup_{\|z\|_2=1} \sqrt{z^T U^T \Sigma^2 U z} = \sup_{\|z\|_2=1} \sqrt{y^T \Sigma^2 y}$$

where $y = Uz$ and U is orthogonal. Since U orthogonal and $\|z\|_2 = 1$, we have $\|y\|_2^2 = \|Uz\|_2^2 = z^T U^T U z = z^T z = \|z\|_2^2 = 1 \implies \|y\|_2 = 1$.

With this property, we have

$$\begin{aligned}
\|A\|_2 &= \sup_{\|y\|_2=1} \sqrt{y^T \Sigma^2 y} = \sup_{\|y\|_2=1} \sqrt{\sum_{i=1}^n \lambda_i^2 y_i^2} \\
&\leq \sqrt{\sum_{i=1}^n \max_i |\lambda_i^2| y_i^2} \\
&= \sqrt{\max_i |\lambda_i^2| \sum_{i=1}^n y_i^2} \\
&= \sqrt{\max_i |\lambda_i^2|} \\
&= \max_i |\lambda_i|
\end{aligned}$$

Assume $\lambda_k = \max_i \lambda_i$, then the "=" is obtained when $y = (0, 0, \dots, 0, \underbrace{1}_{k^{th} \text{ element}}, 0, \dots, 0)^T$ (i.e. only k^{th} element of y is 1 while other elements are 0).

Thus, $\|A\|_2$ = largest of absolute values of eigenvalues of A for symmetric A .

3 Q3

3.1 a

(1) For a rectangular $n \times p$ matrix X , we have SVD of X as $X = U \Sigma V^T =$

$$U \begin{bmatrix} \lambda_1 & & \\ & \ddots & \\ & & \lambda_n \end{bmatrix} V^T \text{ where } V \text{ composed of right singular vectors and } \lambda_i \text{'s are}$$

singular values of X .

Since

$$X^T X = (V \Sigma U^T) U \Sigma V^T = V \Sigma^2 V^T = V \begin{bmatrix} \lambda_1^2 & & \\ & \ddots & \\ & & \lambda_n^2 \end{bmatrix} V^T$$

where V is a matrix of eigenvectors of $X^T X$ and λ_i^2 's are eigenvalues of $X^T X$. Thus, the right singular vectors of X are the eigenvectors of the matrix $X^T X$. Moreover, from above we can see that the eigenvalues of $X^T X$ are the squares of singular values of X .

(2)

$$X^T X = V \Sigma^2 V^T = V \begin{bmatrix} \lambda_1^2 & & \\ & \ddots & \\ & & \lambda_n^2 \end{bmatrix} V^T$$

Since the eigenvalues λ_i^2 's are non-negative, by property of positive semi-definite matrices, we can see that $X^T X$ is positive semi-definite.

3.2 b

Suppose we have eigendecomposition $\Sigma = UAU^T$, where U is a orthonormal matrix and $A = \text{diag}(A_1, \dots, A_n)$ is a diagonal matrix of eigenvalues. Since $U(cI)U^T = c(UIU^T) = c(UU^T) = cI$, we have

$$Z = \Sigma + cI = UAU^T + U(cI)U^T = U(A + cI)U^T = U \begin{bmatrix} A_1 + c & & \\ & \ddots & \\ & & A_n + c \end{bmatrix} U^T$$

Thus, eigenvalues of Z are $(A_1 + c, \dots, A_n + c)$, for which is n additions in total. In this way, we compute $O(n)$ arithmetic calculations.

4 Q4

Collaborate With Ming Qiu

4.1 a

We can first consider the QR decomposition $X = QR$, then we will get $X^T X = R^T R$. Moreover, assume $AR^{-1} = Q_1 R_1$ and will have the following induction:

$$\begin{aligned} \hat{\beta} &= (X^T X)^{-1} X^T Y + (X^T X)^{-1} A^T (A(X^T X)^{-1} A^T)^{-1} (-A(X^T X)^{-1} X^T Y + b) \\ &= (R^T R)^{-1} R^T Q^T Y + (R^T R)^{-1} A^T (A(R^T R)^{-1} A^T)^{-1} (-A(R^T R)^{-1} R^T Q^T Y + b) \\ &= (R^T R)^{-1} R^T Q^T Y + (R^T R)^{-1} A^T (AR^{-1} (AR^{-1})^T)^{-1} (-AR^{-1} Q^T Y + b) \\ &= (R^T R)^{-1} R^T Q^T Y + (R^T R)^{-1} A^T R_1^{-1} (R_1^T)^{-1} (-AR^{-1} Q^T Y + b) \\ &= (R^T R)^{-1} [R^T Q^T Y + A^T R_1^{-1} (R_1^T)^{-1} (-AR^{-1} Q^T Y + b)] \\ &= R^{-1} (R^T)^{-1} [R^T Q^T Y + A^T R_1^{-1} (R_1^T)^{-1} (-AR^{-1} Q^T Y + b)] \end{aligned}$$

For each term in the above calculation with inverse, we use `backsolve()` to get the solutions efficiently.

e.g. get $R^{-1} Q^T Y$ using `backsolve(R, Q^T Y)` and get $(R_1^T)^{-1} (-AR^{-1} Q^T Y + b)$ using `backsolve(R_1, -AR^{-1} Q^T Y + b, transpose=T)`.

4.2 b

In the following code, I implement the above algorithm and showed that it will be more efficient to use QR decomposition and `backsolve()` rather than simply using `solve()` function.

```

get_betahat <- function(A,X,Y,b){
  R <- qr.R(qr(X))
  Q <- qr.Q(qr(X))
  R1 <- qr.R(qr(t(A**solve(R))))
  backsolve(R, backsolve(R, t(X)**Y+t(A) ** backsolve(R1,
    backsolve(R1,-A**backsolve(R, t(Q)**Y)+b, transpose = T )),transpose = T))
}

#Generate random matrix to make some tests
m <- 100
n<- 2000
p <- 2000
set.seed(1)
A <- matrix(rnorm(m*p), nrow = m)
set.seed(1)
X <- matrix(rnorm(n*p), nrow = n)
set.seed(1)
Y <- rnorm(n)
set.seed(1)
b <- rnorm(m)
d <- t(X)**Y

system.time(C <- crossprod(X))

##      user  system elapsed
##    4.50    0.00    4.62

system.time(solve(C)**d +
  solve(C)**t(A)**solve(A**solve(C)**t(A))**(-A**(solve(C)**d)+b))

##      user  system elapsed
##   37.49    0.08   39.50

system.time(get_betahat(A,X,Y,b))

##      user  system elapsed
##   21.61    0.11   22.17

```

From above, we can see that the sum of top 2 elapsed time is much greater than the elapsed time of the third one, which proves my claim.

5 Q5

5.1 a

The reason is that we can't even do the first stage of calculating \hat{X} since although we can calculate $(Z^T Z)^{-1}$ because Z is sparse and $(Z^T Z)$ is only 630×630 , we can't calculate $Z(Z^T Z)^{-1}$ since the result will not be sparse with a extremely huge dimension of 60 million \times 630. When calculating a 60 million \times 630 matrix, it will arise a lack of memory and even though memory is enough, storing such a huge matrix will be another serious issue.

Thus, we can't calculate \hat{X} in this case and can't do the calculation in two stages.

5.2 b

Since $\hat{X} = Z(Z^T Z)^{-1} Z^T X$ and $\hat{\beta} = (\hat{X}^T \hat{X})^{-1} \hat{X}^T y$, we can have the following calculation:

$$\begin{aligned}\hat{\beta} &= (\hat{X}^T \hat{X})^{-1} \hat{X}^T y \\ &= (X^T Z (Z^T Z)^{-1} Z^T Z (Z^T Z)^{-1} Z^T X)^{-1} (Z (Z^T Z)^{-1} Z^T X)^T y \\ &= (X^T Z (Z^T Z)^{-1} Z^T X)^{-1} (X^T Z (Z^T Z)^{-1} Z^T) y \\ &= ((X^T Z) (Z^T Z)^{-1} (Z^T X))^{-1} (X^T Z) (Z^T Z)^{-1} (Z^T y)\end{aligned}$$

Since $Z_{(60 \text{ million} \times 630)}$ and $X_{(60 \text{ million} \times 600)}$ and $y_{(60 \text{ million} \times 1)}$, we can see that the dimensions of pairs in $()$ are $(X^T Z)_{600 \times 630}$, $(Z^T Z)_{630 \times 630}$, $(Z^T X)_{630 \times 600}$ and $(Z^T y)_{630 \times 1}$. In this way, we can actually calculate the values inside the $()$ first (and store them) and then multiply those outputs together (in this case it's really computational non-intensive), which will lead to a applicable solution of $\hat{\beta}$.

6 Q6

In this question, I define error in the estimated eigenvalues relative to the known true values as $\text{error} = \|V_1 - V_2\|_2$, where V_1 = sorted vector of true eigenvalues and V_2 = sorted vector of estimated eigenvalues.

```
set.seed(2)
z <- matrix(rnorm(100*100), nrow = 100)
eigen_vecs <- eigen(crossprod(z))$vectors

#generate a matrix of eigenvalues that are all the same
lambda0 <- diag(x= rep(1.5, 100), nrow=100, ncol=100)

#numerical estimate of above matrix
numerical_lambda0 <- eigen(eigen_vecs %*% lambda0 %*% t(eigen_vecs))$values
```

```

numerical_lambda0 - diag(lambda0)

## [1] 1.150191e-13 4.130030e-14 3.996803e-14 2.198242e-14 1.776357e-14
## [6] 1.421085e-14 1.332268e-14 1.310063e-14 1.132427e-14 1.065814e-14
## [11] 9.769963e-15 9.103829e-15 7.993606e-15 7.327472e-15 6.661338e-15
## [16] 6.439294e-15 5.995204e-15 5.551115e-15 5.551115e-15 5.107026e-15
## [21] 4.662937e-15 4.662937e-15 4.440892e-15 4.440892e-15 4.218847e-15
## [26] 3.996803e-15 3.774758e-15 3.774758e-15 3.552714e-15 3.330669e-15
## [31] 2.886580e-15 2.886580e-15 2.442491e-15 2.442491e-15 2.442491e-15
## [36] 2.220446e-15 2.220446e-15 1.998401e-15 1.998401e-15 1.776357e-15
## [41] 1.554312e-15 1.554312e-15 1.332268e-15 1.332268e-15 1.110223e-15
## [46] 6.661338e-16 6.661338e-16 4.440892e-16 4.440892e-16 2.220446e-16
## [51] 0.000000e+00 -2.220446e-16 -2.220446e-16 -2.220446e-16 -4.440892e-16
## [56] -6.661338e-16 -8.881784e-16 -1.110223e-15 -1.110223e-15 -1.332268e-15
## [61] -1.332268e-15 -1.332268e-15 -1.554312e-15 -1.554312e-15 -1.776357e-15
## [66] -1.776357e-15 -1.998401e-15 -2.220446e-15 -2.442491e-15 -2.664535e-15
## [71] -3.108624e-15 -3.108624e-15 -3.552714e-15 -3.774758e-15 -3.774758e-15
## [76] -3.996803e-15 -4.218847e-15 -4.218847e-15 -4.662937e-15 -4.884981e-15
## [81] -5.329071e-15 -5.329071e-15 -5.551115e-15 -6.217249e-15 -6.439294e-15
## [86] -7.105427e-15 -7.327472e-15 -8.881784e-15 -9.547918e-15 -9.769963e-15
## [91] -1.021405e-14 -1.154632e-14 -1.287859e-14 -1.332268e-14 -1.554312e-14
## [96] -1.820766e-14 -2.109424e-14 -3.907985e-14 -4.041212e-14 -1.130207e-13

numerical_lambda0

## [1] 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5
## [18] 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5
## [35] 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5
## [52] 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5
## [69] 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5
## [86] 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5 1.5

#generate matrices of eigenvalues vary between range of values from
#very large to very small, and find the magnitude of condition number
#that the matrix is not numerically PSD
i=1
min=1
cond_num= c()
error=c()
while (min >0){
  set.seed(2)
  # increase the magnitude of eigenvalues and generate a random vector
  # of eigenvalues for each magnitude
  eigen_vals = c(10^i, sort(runif(98, min =1e-4, max =10^i),
    decreasing= TRUE), 1e-4)
  lambda1 <- diag(x = eigen_vals, nrow=100,ncol=100)

```

```

num_lbda1 <- eigen(eigen_vecs %*% lambda1 %*% t(eigen_vecs))$values
num_lbda1 <- sort(num_lbda1, decreasing = TRUE)

error <- c(error, sum((eigen_vals - num_lbda1)^2) )
cond_num<- c(cond_num, max(num_lbda1)/min(num_lbda1))
min= min(num_lbda1)
i= i+1
}

# Print Error Values
error

## [1] 1.497919e-27 1.733416e-25 1.425427e-23 1.172706e-21 1.959432e-19
## [6] 1.470017e-17 1.463941e-15 1.483236e-13 1.431569e-11 1.240705e-09
## [11] 1.414853e-07 1.294964e-05 1.355427e-03 1.222095e-01 8.214544e+00
## [16] 9.428059e+02 9.695232e+04 8.454407e+06

# Print Condition Numbers
cond_num

## [1] 1.000000e+05 1.000000e+06 1.000000e+07 1.000000e+08 9.999997e+08
## [6] 9.999947e+09 9.999644e+10 9.996851e+11 1.002223e+13 9.892226e+13
## [11] 7.281778e+14 4.096000e+15 2.560000e+15 2.124985e+15 1.472002e+16
## [16] 1.034188e+16 5.549190e+15 -9.951082e+15

# emperical conditional number for the matrix that is not numerically psd
abs(cond_num[i-1])

## [1] 9.951082e+15

# True condition number
10^(i+4-1)

## [1] 1e+22

```

When emperical condition number is at 9.951082e+15 magnitude and true condition number is at 1e+22 magnitude, the matrix is not numerically positive semi-definite.

Moreover, from above values of errors and condition numbers, we can see that (1)the error approximately increases in 10^2 magnitude as the eigenvalues increase in 10^1 magnitude, and (2) the error increases as condition number increases until condition number become negative(at which point the matrix is no more numerically positive semi-definite).